

Glass Boxes over Black Boxes

An Empirical Study on Fairness in Criminal Justice AI

Authors: Vidhi Bhatt, Yuri Kim, Ashvin Loghashankar

Course: CS 269 Fall 2025, UCLA

Project Overview

This study compares interpretable models (Decision Trees, Logistic Regression) against black-box models (Random Forest, XGBoost) on the ProPublica COMPAS dataset, focusing on:

1. **Performance Comparison:** Accuracy, precision, recall, and F1-scores
2. **Fairness Analysis:** Disparate impact and error rate differences across racial groups
3. **Bias Mitigation:** Feature removal and fairness-constrained training
4. **Causal Analysis:** Decomposing bias into data-driven vs algorithmic components

In [1]:

```
# Installing required packages for fairness analysis
!pip install fairlearn
!pip install shap
!pip install xgboost

Collecting fairlearn
  Downloading fairlearn-0.13.0-py3-none-any.whl.metadata (7.3 kB)
Requirement already satisfied: narwhals>=1.14.0 in /usr/local/lib/python3.12/dist-packages (from fairlearn) (2.12.0)
Requirement already satisfied: numpy>=1.24.4 in /usr/local/lib/python3.12/dist-packages (from fairlearn) (2.0.2)
Requirement already satisfied: pandas>=2.0.3 in /usr/local/lib/python3.12/dist-packages (from fairlearn) (2.2.2)
Requirement already satisfied: scikit-learn>=1.2.1 in /usr/local/lib/python3.12/dist-packages (from fairlearn) (1.6.1)
Collecting scipy<1.16.0,>=1.9.3 (from fairlearn)
  Downloading scipy-1.15.3-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (61 kB)
                                                 62.0/62.0 kB 2.4 MB/s eta 0:00:00
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=2.0.3->fairlearn) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=2.0.3->fairlearn) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=2.0.3->fairlearn) (2025.2)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn>=1.2.1->fairlearn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn>=1.2.1->fairlearn) (3.6.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas>=2.0.3->fairlearn) (1.17.0)
  Downloading fairlearn-0.13.0-py3-none-any.whl (251 kB)
                                                 251.6/251.6 kB 2.7 MB/s eta 0:00:00
Downloaded scipy-1.15.3-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (37.3 MB)
                                                 37.3/37.3 MB 22.0 MB/s eta 0:00:00
Installing collected packages: scipy, fairlearn
  Attempting uninstall: scipy
    Found existing installation: scipy 1.16.3
    Uninstalling scipy-1.16.3:
      Successfully uninstalled scipy-1.16.3
  Successfully installed fairlearn-0.13.0 scipy-1.15.3
```

```
Requirement already satisfied: shap in /usr/local/lib/python3.12/dist-packages (0.50.0)
Requirement already satisfied: numpy>=2 in /usr/local/lib/python3.12/dist-packages (from shap) (2.0.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from shap) (1.15.3)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (from shap) (1.6.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (from shap) (2.2.2)
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.12/dist-packages (from shap) (4.67.1)
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.12/dist-packages (from shap) (25.0)
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.12/dist-packages (from shap) (0.0.8)
Requirement already satisfied: numba>=0.54 in /usr/local/lib/python3.12/dist-packages (from shap) (0.60.0)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.12/dist-packages (from shap) (3.1.2)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.12/dist-packages (from shap) (4.15.0)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.12/dist-packages (from numba>=0.54->shap) (0.43.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas->shap) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas->shap) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas->shap) (2025.2)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn->shap) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn->shap) (3.6.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas->shap) (1.17.0)
Requirement already satisfied: xgboost in /usr/local/lib/python3.12/dist-packages (3.1.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from xgboost) (2.0.2)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.12/dist-packages (from xgboost) (2.27.5)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from xgboost) (1.15.3)
```

In [2]:

```
# Importing essential libraries for data analysis and machine learning
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Importing scikit-Learn modules for model training and evaluation
from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score

# Importing fairness Libraries for bias mitigation
from fairlearn.reductions import GridSearch, EqualizedOdds, DemographicParity
from fairlearn.postprocessing import ThresholdOptimizer
from fairlearn.metrics import demographic_parity_difference, equalized_odds_difference

# Importing statistical analysis tools
from scipy.stats import chi2_contingency, pointbiserialr, ttest_ind
```

```
import warnings
warnings.filterwarnings('ignore')

# Setting up visualization parameters
sns.set_style('whitegrid')
plt.rcParams['figure.figsize'] = (14, 7)
```

In [4]:

```
# Loading the ProPublica COMPAS dataset for fairness analysis
df = pd.read_csv('propublica_data_for_fairml.csv')
```

In [5]:

```
# Performing basic exploratory data analysis
print(df.describe())

# Examining the distribution of our target variable (recidivism)
print(df['Two_yr_Recidivism'].value_counts())
print(f"\nRecidivism Rate: {df['Two_yr_Recidivism'].mean():.2%}")
```

	Two_yr_Recidivism	Number_of_Priors	score_factor	\
count	6172.000000	6172.000000	6172.000000	
mean	0.455120	3.246436	0.445723	
std	0.498022	4.743770	0.497086	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	
50%	0.000000	1.000000	0.000000	
75%	1.000000	4.000000	1.000000	
max	1.000000	38.000000	1.000000	

	Age_Above_FourtyFive	Age_Below_TwentyFive	African_American	\
count	6172.000000	6172.000000	6172.000000	
mean	0.209494	0.218244	0.514420	
std	0.406981	0.413087	0.499833	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	1.000000	
75%	0.000000	0.000000	1.000000	
max	1.000000	1.000000	1.000000	

	Asian	Hispanic	Native_American	Other	Female	\
count	6172.000000	6172.000000	6172.000000	6172.000000	6172.000000	
mean	0.005023	0.082469	0.001782	0.055574	0.190376	
std	0.070698	0.275101	0.042182	0.229115	0.392629	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	Misdemeanor					
count	6172.000000					
mean	0.356773					
std	0.479086					
min	0.000000					
25%	0.000000					
50%	0.000000					
75%	1.000000					
max	1.000000					

	Two_yr_Recidivism					
0	3363					
1	2809					

Name: count, dtype: int64

Recidivism Rate: 45.51%

In [6]:

```
# Analyzing demographic distribution in the dataset
demographic_cols = ['African_American', 'Asian', 'Hispanic', 'Native_American', 'Oth
```

```

for col in demographic_cols:
    print(f'{col}: {df[col].sum()} ({df[col].mean():.2%})')

# Creating a consolidated race variable for easier analysis
df['race'] = 'Other'
df.loc[df['African_American'] == 1, 'race'] = 'African American'
df.loc[df['Asian'] == 1, 'race'] = 'Asian'
df.loc[df['Hispanic'] == 1, 'race'] = 'Hispanic'
df.loc[df['Native_American'] == 1, 'race'] = 'Native American'

print("\nRace Distribution from data:")
print(df['race'].value_counts())

```

African_American: 3175 (51.44%)
 Asian: 31 (0.50%)
 Hispanic: 509 (8.25%)
 Native_American: 11 (0.18%)
 Other: 343 (5.56%)
 Female: 1175 (19.04%)

Race Distribution from data:

race	count
African American	3175
Other	2446
Hispanic	509
Asian	31
Native American	11
Name: count, dtype: int64	

In [7]:

```

# Calculating recidivism rates by racial group to identify disparities
recidivism_by_race = df.groupby('race')['Two_yr_Recidivism'].agg(['sum', 'count', 'mean'])
recidivism_by_race.columns = ['Recidivists', 'Total', 'Recidivism_Rate']
print("Recidivism by Race:")
print(recidivism_by_race)

# Visualizing racial disparities in recidivism rates
fig, ax = plt.subplots(1, 2, figsize=(14, 5))

recidivism_by_race['Recidivism_Rate'].plot(kind='bar', ax=ax[0], color='steelblue')
ax[0].set_title('Recidivism Rate by Race', fontsize=12, fontweight='bold')
ax[0].set_ylabel('Recidivism Rate')
ax[0].set_xlabel('Race')
ax[0].set_ylim([0, 1])
ax[0].axhline(y=df['Two_yr_Recidivism'].mean(), color='red', linestyle='--', label='Mean')
ax[0].legend()
plt.setp(ax[0].xaxis.get_majorticklabels(), rotation=45, ha='right')

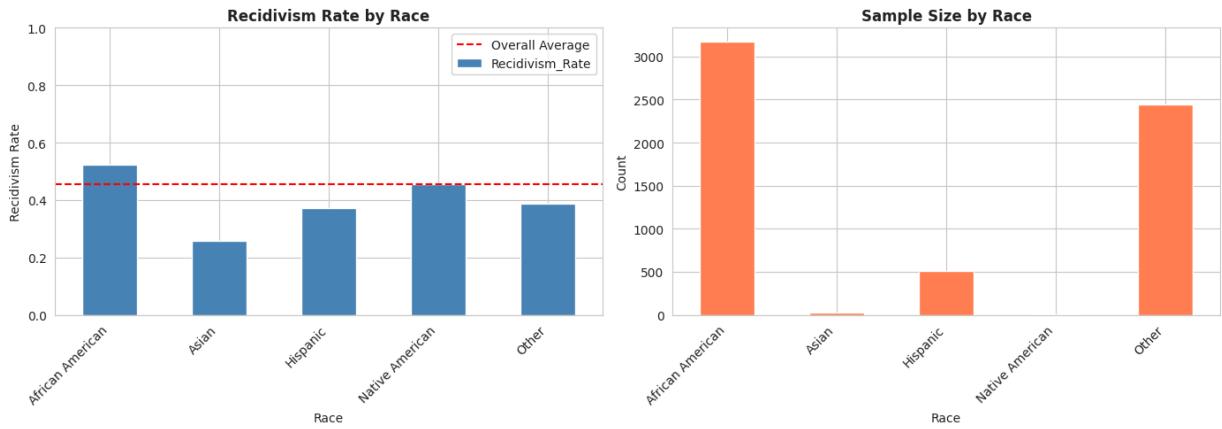
recidivism_by_race['Total'].plot(kind='bar', ax=ax[1], color='coral')
ax[1].set_title('Sample Size by Race', fontsize=12, fontweight='bold')
ax[1].set_ylabel('Count')
ax[1].set_xlabel('Race')
plt.setp(ax[1].xaxis.get_majorticklabels(), rotation=45, ha='right')

plt.tight_layout()
plt.show()

```

Recidivism by Race:

race	Recidivists	Total	Recidivism_Rate
African American	1661	3175	0.523150
Asian	8	31	0.258065
Hispanic	189	509	0.371316
Native American	5	11	0.454545
Other	946	2446	0.386754



In [8]:

```
# Preparing features and target variables for machine Learning
X = df.drop(['Two_yr_Recidivism', 'race'], axis=1)
y = df['Two_yr_Recidivism']
sensitive_attr = df['African_American'] # 1 = African American, 0 = Other

# Splitting data while maintaining class balance
X_train, X_test, y_train, y_test, sens_train, sens_test = train_test_split(
    X, y, sensitive_attr, test_size=0.3, random_state=42, stratify=y
)
```

In [9]:

```
# Standardizing features for algorithms sensitive to scale
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

In [10]:

```
# Training our first interpretable model - Decision Tree
dt_model = DecisionTreeClassifier(max_depth=5, random_state=42, min_samples_split=20)
dt_model.fit(X_train, y_train)
# Generating predictions for evaluation
dt_pred = dt_model.predict(X_test)
dt_pred_proba = dt_model.predict_proba(X_test)[:, 1]

# Examining model complexity
print(f"Tree Depth: {dt_model.get_depth()}")
print(f"Number of Leaves: {dt_model.get_n_leaves()}")
```

Tree Depth: 5
Number of Leaves: 30

In [11]:

```
# Training our second interpretable model - Logistic Regression
lr_model = LogisticRegression(max_iter=1000, random_state=42)
lr_model.fit(X_train_scaled, y_train)
# Generating predictions for evaluation
lr_pred = lr_model.predict(X_test_scaled)
lr_pred_proba = lr_model.predict_proba(X_test_scaled)[:, 1]

# Examining feature coefficients to understand model decisions
coef_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': lr_model.coef_[0]
}).sort_values('Coefficient', ascending=False)
print(coef_df)
```

	Feature	Coefficient
0	Number_of_Priors	0.596545
1	score_factor	0.364085

```

3   Age_Below_TwentyFive      0.206868
8       Other                 0.008456
7   Native_American          -0.005390
4   African_American         -0.012557
5       Asian                -0.048261
6   Hispanic                -0.067755
10  Misdemeanor             -0.131940
9       Female               -0.157924
2   Age_Above_FourtyFive    -0.190544

```

In [12]:

```

# Training our first black box model - Random Forest
rf_model = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42, n_jobs=-1)
rf_model.fit(X_train, y_train)
# Generating predictions for evaluation
rf_pred = rf_model.predict(X_test)
rf_pred_proba = rf_model.predict_proba(X_test)[:, 1]

```

In [13]:

```

# Training our second black box model - XGBoost
xgb_model = XGBClassifier(n_estimators=100, max_depth=5, random_state=42, verbosity=0)
xgb_model.fit(X_train, y_train)
# Generating predictions for evaluation
xgb_pred = xgb_model.predict(X_test)
xgb_pred_proba = xgb_model.predict_proba(X_test)[:, 1]

```

In [14]:

```

# Computing standard performance metrics for all models
def compute_metrics(y_true, y_pred, model_name):
    accuracy = accuracy_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred)
    recall = recall_score(y_true, y_pred)

    return {
        'Model': model_name,
        'Accuracy': accuracy,
        'F1-Score': f1,
        'Precision': precision,
        'Recall': recall
    }

# Evaluating all trained models
metrics_list = [
    compute_metrics(y_test, dt_pred, 'Decision Tree'),
    compute_metrics(y_test, lr_pred, 'Logistic Regression'),
    compute_metrics(y_test, rf_pred, 'Random Forest'),
    compute_metrics(y_test, xgb_pred, 'XGBoost')
]

metrics_df = pd.DataFrame(metrics_list)
print("\nAccuracy metrics:")
print(metrics_df.to_string(index=False))

```

Accuracy metrics:

	Model	Accuracy	F1-Score	Precision	Recall
Decision Tree	0.677646	0.621433	0.667575	0.581257	
Logistic Regression	0.684125	0.623794	0.681180	0.575326	
Random Forest	0.672246	0.614113	0.661644	0.572954	
XGBoost	0.675486	0.609487	0.673851	0.556346	

Fairness Metrics

This section evaluates algorithmic bias by measuring disparate impact ratios and error rate differences across racial groups. We focus on three key metrics:

- **Disparate Impact:** Ratio of positive prediction rates between African Americans and other races
- **False Positive Rate (FPR) Difference:** Difference in incorrectly flagging low-risk individuals
- **False Negative Rate (FNR) Difference:** Difference in missing high-risk individuals

In [15]:

```
# Computing fairness metrics to assess algorithmic bias
# sensitive_attr: 1 = African American, 0 = Other
def compute_fairness_metrics(y_true, y_pred, sensitive_attr, model_name):
    # Calculating ratio of positive predictions for African Americans vs Others
    aa_positive_rate = y_pred[sensitive_attr == 1].mean()
    other_positive_rate = y_pred[sensitive_attr == 0].mean()
    disparate_impact = aa_positive_rate / (other_positive_rate + 1e-6)

    # Computing False Positive Rate differences
    aa_fpr = false_positive_rate(y_true[sensitive_attr == 1], y_pred[sensitive_attr == 1])
    other_fpr = false_positive_rate(y_true[sensitive_attr == 0], y_pred[sensitive_attr == 0])
    fpr_difference = aa_fpr - other_fpr

    # Computing False Negative Rate differences
    aa_fnr = false_negative_rate(y_true[sensitive_attr == 1], y_pred[sensitive_attr == 1])
    other_fnr = false_negative_rate(y_true[sensitive_attr == 0], y_pred[sensitive_attr == 0])
    fnr_difference = aa_fnr - other_fnr

    return {
        'Model': model_name,
        'AA_Positive_Rate': aa_positive_rate,
        'Other_Positive_Rate': other_positive_rate,
        'Disparate_Impact_Ratio': disparate_impact,
        'AA_FPR': aa_fpr,
        'Other_FPR': other_fpr,
        'FPR_Difference': fpr_difference,
        'AA_FNR': aa_fnr,
        'Other_FNR': other_fnr,
        'FNR_Difference': fnr_difference
    }

# Evaluating fairness across all trained models
fairness_list = [
    compute_fairness_metrics(y_test, dt_pred, sens_test, 'Decision Tree'),
    compute_fairness_metrics(y_test, lr_pred, sens_test, 'Logistic Regression'),
    compute_fairness_metrics(y_test, rf_pred, sens_test, 'Random Forest'),
    compute_fairness_metrics(y_test, xgb_pred, sens_test, 'XGBoost')
]

fairness_df = pd.DataFrame(fairness_list)
print("\nFairness Metrics:")
print(fairness_df.to_string(index=False))
```

Fairness Metrics:

	Model	AA_Positive_Rate	Other_Positive_Rate	Disparate_Impact_Ratio			
	AA_FPR	Other_FPR	FPR_Difference	AA_FNR	Other_FNR	FNR_Difference	
Decision Tree		0.512552		0.272321			1.882152 0.
339326	0.164894	0.174432	0.336595	0.545181		-0.208586	
Logistic Regression		0.517782		0.242188			2.137932 0.
330337	0.141844	0.188493	0.318982	0.587349		-0.268367	
Random Forest		0.521967		0.257812			2.024590 0.
355056	0.157801	0.197255	0.332681	0.572289		-0.239608	
XGBoost		0.491632		0.252232			1.949117 0.
325843	0.145390	0.180453	0.363992	0.566265		-0.202273	

In [16]:

```
# Comparing interpretable vs black-box models on accuracy and fairness
print(f"\nInterpretable Models Average Accuracy: {metrics_df[metrics_df['Model'].isin(['Logistic Regression', 'Support Vector Machine'])].mean()['Accuracy']:.4f}")
print(f"Black-Box Models Average Accuracy: {metrics_df[metrics_df['Model'].isin(['Random Forest', 'Gradient Boosting', 'XGBoost'])].mean()['Accuracy']:.4f}")
print(f"\nInterpretable Models Average Disparate Impact: {fairness_df[fairness_df['Model'].isin(['Logistic Regression', 'Support Vector Machine'])].mean()['Disparate Impact']:.4f}")
print(f"Black-Box Models Average Disparate Impact: {fairness_df[fairness_df['Model'].isin(['Random Forest', 'Gradient Boosting', 'XGBoost'])].mean()['Disparate Impact']:.4f}")
```

Interpretable Models Average Accuracy: 0.6809

Black-Box Models Average Accuracy: 0.6739

Interpretable Models Average Disparate Impact: 2.0100

Black-Box Models Average Disparate Impact: 1.9869

In [17]:

```
# Conducting ablation study to understand feature impact on fairness
def fairness_metric(y_true, y_pred, sensitive_attr):
    aa_positive_rate = y_pred[sensitive_attr == 1].mean()
    other_positive_rate = y_pred[sensitive_attr == 0].mean()
    disparate_impact = aa_positive_rate / (other_positive_rate + 1e-6)
    return disparate_impact

ablation_results = []

# Establishing baseline performance with all features
lr_baseline = LogisticRegression(max_iter=1000, random_state=42)
lr_baseline.fit(X_train_scaled, y_train)
baseline_pred = lr_baseline.predict(X_test_scaled)
baseline_di = fairness_metric(y_test, baseline_pred, sens_test)
baseline_acc = accuracy_score(y_test, baseline_pred)

ablation_results.append({
    'Configuration': 'Baseline (All Features)',
    'Features_Removed': 'None',
    'Num_Features': X_train.shape[1],
    'Disparate_Impact': baseline_di,
    'Accuracy': baseline_acc,
    'DI_Improvement': 0,
    'Accuracy_Loss': 0
})
print("Impact of Removing Features on Disparate Impact")
print(f"\nBaseline (All Features):")
print(f"  Disparate Impact: {baseline_di:.4f}")
print(f"  Accuracy: {baseline_acc:.4f}")

# Testing impact of removing each feature individually
for feature_to_remove in X_train.columns:
    X_train_removed = X_train.drop(feature_to_remove, axis=1)
    X_test_removed = X_test.drop(feature_to_remove, axis=1)

    # Retraining model without the selected feature
    scaler_temp = StandardScaler()
    X_train_temp = scaler_temp.fit_transform(X_train_removed)
    X_test_temp = scaler_temp.transform(X_test_removed)

    lr_temp = LogisticRegression(max_iter=1000, random_state=42)
    lr_temp.fit(X_train_temp, y_train)
    temp_pred = lr_temp.predict(X_test_temp)

    # Measuring changes in fairness and accuracy
    temp_di = fairness_metric(y_test, temp_pred, sens_test)
    temp_acc = accuracy_score(y_test, temp_pred)

    di_improvement = baseline_di - temp_di
    acc_loss = baseline_acc - temp_acc
```

```

ablation_results.append({
    'Configuration': f'Remove: {feature_to_remove}',
    'Features_Removed': feature_to_remove,
    'Num_Features': X_train_removed.shape[1],
    'Disparate_Impact': temp_di,
    'Accuracy': temp_acc,
    'DI_Improvement': di_improvement,
    'Accuracy_Loss': acc_loss
})

ablation_df = pd.DataFrame(ablation_results).sort_values('DI_Improvement', ascending=False)
print(ablation_df.to_string(index=False))

```

Impact of Removing Features on Disparate Impact

Baseline (All Features):

Disparate Impact: 2.1379

Accuracy: 0.6841

	Configuration	Features_Removed	Num_Features	Disparate_Impact	Accuracy
	DI_Improvement	Accuracy_Loss			
663607	Remove: Number_of_Priors 0.105943	Number_of_Priors 0.020518	10	2.031988	0.
683585	Remove: Female 0.038526	Female 0.000540	10	2.099406	0.
683585	Remove: Native_American 0.009807	Native_American 0.000540	10	2.128125	0.
684125	Baseline (All Features) 0.000000	None 0.000000	11	2.137932	0.
684125	Remove: Other 0.000000	Other 0.000000	10	2.137932	0.
677646	Remove: Misdemeanor -0.040271	Misdemeanor 0.006479	10	2.178203	0.
679806	Remove: Age_Above_FourtyFive -0.066801	Age_Above_FourtyFive 0.004320	10	2.204733	0.
684125	Remove: Asian -0.081834	Asian 0.000000	10	2.219766	0.
681965	Remove: Hispanic -0.097012	Hispanic 0.002160	10	2.234944	0.
671166	Remove: Age_Below_TwentyFive -0.111431	Age_Below_TwentyFive 0.012959	10	2.249363	0.
683045	Remove: African_American -0.112337	African_American 0.001080	10	2.250268	0.
680346	Remove: score_factor -0.278896	score_factor 0.003780	10	2.416827	0.

In [18]:

```

# Visualizing the fairness-accuracy trade-offs from ablation study
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

ablation_sorted = ablation_df[ablation_df['Configuration'] != 'Baseline (All Feature']
colors = ['green' if x < baseline_di else 'red' for x in ablation_sorted['Disparate_Impact']]
axes[0].barh(ablation_sorted['Configuration'], ablation_sorted['Disparate_Impact'],
            axes[0].axvline(x=baseline_di, color='black', linestyle='--', linewidth=2, label=f'B'),
            axes[0].axvline(x=1.0, color='green', linestyle=':', linewidth=2, label='Perfect Fairness'),
            axes[0].set_xlabel('Disparate Impact Ratio', fontweight='bold'),
            axes[0].set_title('Disparate Impact: Impact of Feature Removal', fontweight='bold'),
            axes[0].legend(),
            axes[0].set_xlim([1, 2.5])

ablation_for_plot = ablation_df[ablation_df['Configuration'] != 'Baseline (All Features)']
scatter = axes[1].scatter(ablation_for_plot['Accuracy_Loss'],
                         ablation_for_plot['DI_Improvement'],
                         s=200, alpha=0.6, c=range(len(ablation_for_plot)),
                         cmap='viridis', edgecolors='black', linewidth=1.5)
axes[1].scatter([0], [0], s=300, marker='*', c='red', edgecolors='black', linewidth=2)
axes[1].axhline(y=0, color='gray', linestyle='-', alpha=0.3)

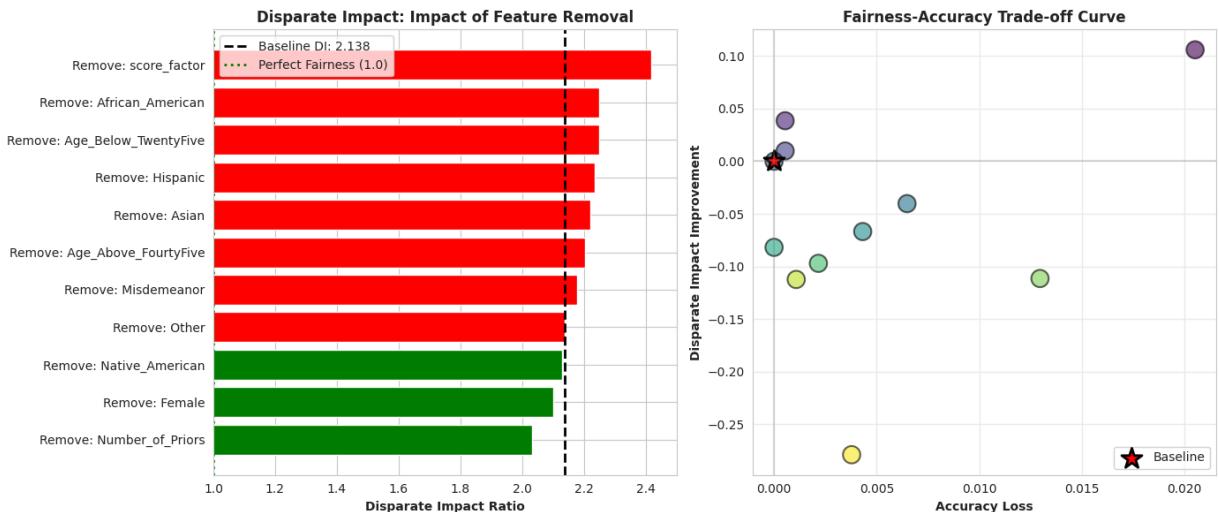
```

```

    axes[1].axvline(x=0, color='gray', linestyle='-', alpha=0.3)
    axes[1].set_xlabel('Accuracy Loss', fontweight='bold')
    axes[1].set_ylabel('Disparate Impact Improvement', fontweight='bold')
    axes[1].set_title('Fairness-Accuracy Trade-off Curve', fontweight='bold')
    axes[1].legend()
    axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



In [19]:

```

# Implementing fairness-constrained training approaches
print("Fairness-constrained training")

# Applying Demographic Parity constraint (Equal positive rates across groups)
print("\n1. DEMOGRAPHIC PARITY CONSTRAINT")
print("    Goal: Equal prediction rates across racial groups")
print("    Constraint: P(Y_pred=1|AA=1) = P(Y_pred=1|AA=0)")

# Training model with demographic parity constraint
constraint_dp = DemographicParity(difference_bound=0.1)
mitigator_dp = GridSearch(
    LogisticRegression(max_iter=1000, random_state=42),
    constraints=constraint_dp,
    grid_size=50
)

mitigator_dp.fit(X_train_scaled, y_train, sensitive_features=sens_train)
dp_pred = mitigator_dp.predict(X_test_scaled)

# Evaluating constrained model performance
dp_accuracy = accuracy_score(y_test, dp_pred)
dp_di = fairness_metric(y_test, dp_pred, sens_test)

print(f"\nResults with Demographic Parity:")
print(f"    Accuracy: {dp_accuracy:.4f} (vs baseline {baseline_acc:.4f}, change: {dp_a})
print(f"    Disparate Impact: {dp_di:.4f} (vs baseline {baseline_di:.4f}, improvement: {dp_improvement:.4f}")
dp_applied = True

```

Fairness-constrained training

1. DEMOGRAPHIC PARITY CONSTRAINT

Goal: Equal prediction rates across racial groups
 Constraint: $P(Y_{pred}=1|AA=1) = P(Y_{pred}=1|AA=0)$

Results with Demographic Parity:

Accuracy: 0.6782 (vs baseline 0.6841, change: -0.0059)
 Disparate Impact: 1.1776 (vs baseline 2.1379, improvement: 0.9603)

In [20]:

```

# Applying Equalized Odds constraint (Equal error rates across groups)
print("\n2. Equalized Odds constraint")
print("    Goal: Equal false positive and false negative rates across groups")
print("    Constraints: FPR_AA = FPR_Other AND FNR_AA = FNR_Other")

# Training model with equalized odds constraint
constraint_eo = EqualizedOdds(difference_bound=0.1)
mitigator_eo = GridSearch(
    LogisticRegression(max_iter=1000, random_state=42),
    constraints=constraint_eo,
    grid_size=50
)

mitigator_eo.fit(X_train_scaled, y_train, sensitive_features=sens_train)
eo_pred = mitigator_eo.predict(X_test_scaled)

# Evaluating equalized odds model performance
eo_accuracy = accuracy_score(y_test, eo_pred)
eo_di = fairness_metric(y_test, eo_pred, sens_test)

print(f"\nResults with Equalized Odds:")
print(f"    Accuracy: {eo_accuracy:.4f} (vs baseline {baseline_acc:.4f}, change: {eo_a})
print(f"    Disparate Impact: {eo_di:.4f} (vs baseline {baseline_di:.4f}, improvement:
eo_applied = True

```

2. Equalized Odds constraint

Goal: Equal false positive and false negative rates across groups
 Constraints: FPR_AA = FPR_Other AND FNR_AA = FNR_Other

Results with Equalized Odds:

Accuracy: 0.6582 (vs baseline 0.6841, change: -0.0259)
 Disparate Impact: 1.0879 (vs baseline 2.1379, improvement: 1.0501)

In [21]:

```

# Comparing effectiveness of different fairness interventions
intervention_results = [
    {
        'Intervention': 'Baseline (No Constraints)',
        'Accuracy': baseline_acc,
        'Disparate_Impact': baseline_di,
        'Accuracy_Loss': 0,
        'Fairness_Gain': 0,
        'Efficiency_Ratio': np.inf
    }
]

# Recording demographic parity results if applied
if dp_applied and dp_accuracy is not None:
    intervention_results.append({
        'Intervention': 'Demographic Parity Constraint',
        'Accuracy': dp_accuracy,
        'Disparate_Impact': dp_di,
        'Accuracy_Loss': baseline_acc - dp_accuracy,
        'Fairness_Gain': baseline_di - dp_di,
        'Efficiency_Ratio': (baseline_di - dp_di) / max(baseline_acc - dp_accuracy,
    })

# Recording equalized odds results if applied
if eo_applied and eo_accuracy is not None:
    intervention_results.append({
        'Intervention': 'Equalized Odds Constraint',
        'Accuracy': eo_accuracy,
        'Disparate_Impact': eo_di,
        'Accuracy_Loss': baseline_acc - eo_accuracy,
    })

```

```

        'Fairness_Gain': baseline_di - eo_di,
        'Efficiency_Ratio': (baseline_di - eo_di) / max(baseline_acc - eo_accuracy,
    })

intervention_df = pd.DataFrame(intervention_results)
print("\n" + intervention_df.to_string(index=False))

```

	Intervention	Accuracy	Disparate_Impact	Accuracy_Loss	Fairness_Gain
00	Baseline (No Constraints)	0.684125	2.137932	0.000000	0.0000
01	Demographic Parity Constraint	0.678186	1.177631	0.005940	0.9603
02	Equalized Odds Constraint	0.658207	1.087863	0.025918	1.0500
68		40.515136			

In [22]:

```

# Plotting
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

intervention_plot = intervention_df[intervention_df['Intervention'] != 'Baseline (No Constraints)']
scatter = axes[0].scatter(intervention_plot['Accuracy_Loss'],
                         intervention_plot['Fairness_Gain'],
                         s=300, alpha=0.7, c=range(len(intervention_plot)),
                         cmap='viridis', edgecolors='black', linewidth=2)
axes[0].scatter([0], [0], s=400, marker='*', c='red', edgecolors='black', linewidth=2)

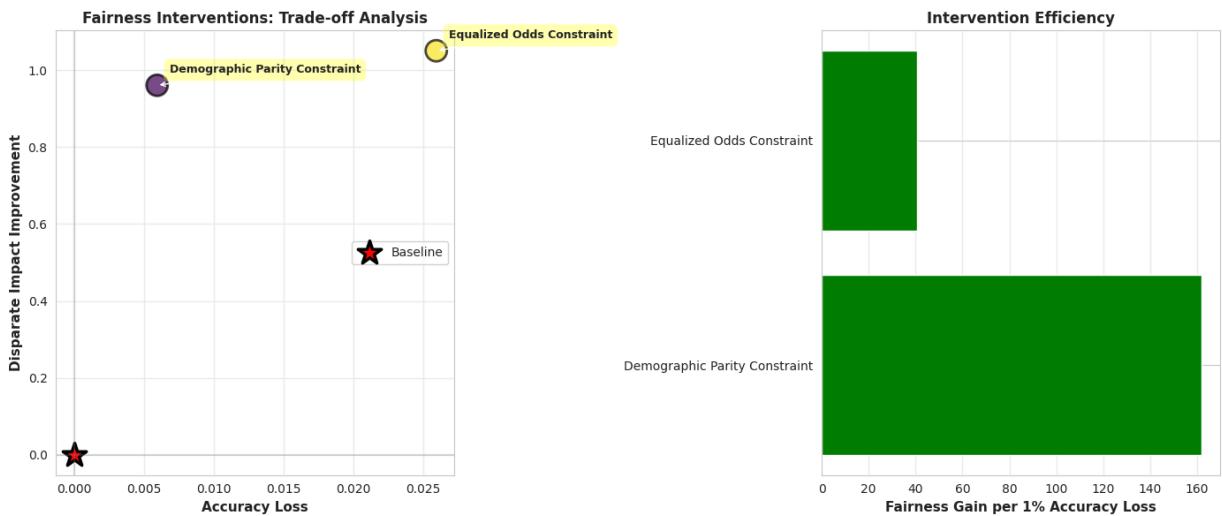
for idx, row in intervention_plot.iterrows():
    axes[0].annotate(row['Intervention'],
                     (row['Accuracy_Loss'], row['Fairness_Gain']),
                     fontsize=9, fontweight='bold',
                     xytext=(10, 10), textcoords='offset points',
                     bbox=dict(boxstyle='round,pad=0.5', facecolor='yellow', alpha=0.3),
                     arrowprops=dict(arrowstyle='->', connectionstyle='arc3,rad=0'))

axes[0].axhline(y=0, color='gray', linestyle='-', alpha=0.3)
axes[0].axvline(x=0, color='gray', linestyle='-', alpha=0.3)
axes[0].set_xlabel('Accuracy Loss', fontweight='bold', fontsize=11)
axes[0].set_ylabel('Disparate Impact Improvement', fontweight='bold', fontsize=11)
axes[0].set_title('Fairness Interventions: Trade-off Analysis', fontweight='bold', fontsize=14)
axes[0].legend()
axes[0].grid(True, alpha=0.3)

intervention_eff = intervention_df[intervention_df['Efficiency_Ratio'] != np.inf].sort_values(['Efficiency_Ratio'])
colors_eff = ['green' if x > 0 else 'gray' for x in intervention_eff['Efficiency_Ratio']]
axes[1].barh(intervention_eff['Intervention'], intervention_eff['Efficiency_Ratio'],
            color=colors_eff)
axes[1].set_xlabel('Fairness Gain per 1% Accuracy Loss', fontweight='bold', fontsize=11)
axes[1].set_title('Intervention Efficiency', fontweight='bold', fontsize=12)
axes[1].grid(True, alpha=0.3, axis='x')

plt.tight_layout()
plt.show()

```



Statistical Significance Testing

We use cross-validation and t-tests to determine whether observed performance differences between interpretable and black-box models are statistically meaningful or due to random variation.

In [23]:

```
# Performing cross-validation to assess model reliability
models_cv = {
    'Logistic Regression': LogisticRegression(max_iter=1000, random_state=42),
    'Decision Tree': DecisionTreeClassifier(max_depth=5, random_state=42, min_samples_leaf=1)
}

# Using stratified k-fold to maintain class balance across folds
kfolds = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
cv_results = {}

# Evaluating each model across multiple folds
for name, model in models_cv.items():
    scores = cross_val_score(model, X_train_scaled, y_train, cv=kfolds, scoring='accuracy')
    cv_results[name] = scores
    print(f"\n{name}:")
    print(f"  Fold Scores: {[f'{s:.4f}' for s in scores]}")
    print(f"  Mean: {scores.mean():.4f}")
    print(f"  Std Dev: {scores.std():.4f}")
    print(f"  95% CI: [{scores.mean() - 1.96*scores.std()/np.sqrt(5):.4f}, {scores.mean() + 1.96*scores.std()/np.sqrt(5):.4f}]")
```

Logistic Regression:

```
Fold Scores: ['0.6910', '0.6944', '0.6562', '0.6667', '0.6655']
Mean: 0.6748
Std Dev: 0.0151
95% CI: [0.6615, 0.6880]
```

Decision Tree:

```
Fold Scores: ['0.7095', '0.7130', '0.6736', '0.6933', '0.6690']
Mean: 0.6917
Std Dev: 0.0180
95% CI: [0.6759, 0.7074]
```

In [24]:

```
# Conducting two-sample t-tests to compare model performance

model_names = list(cv_results.keys())
for i in range(len(model_names)):
    for j in range(i+1, len(model_names)):
        name1, name2 = model_names[i], model_names[j]
```

```
scores1, scores2 = cv_results[name1], cv_results[name2]

# Performing statistical test for significant differences
t_stat, p_value = ttest_ind(scores1, scores2)

print(f"\n{name1} vs {name2}:")
print(f" Mean Difference: {scores1.mean() - scores2.mean():.4f}")
print(f" t-statistic: {t_stat:.4f}")
print(f" p-value: {p_value:.4f}")

# Interpreting statistical significance
if p_value < 0.05:
    print(f" Result: STATISTICALLY SIGNIFICANT (p < 0.05)")
    if scores1.mean() > scores2.mean():
        print(f" - {name1} is significantly better")
    else:
        print(f" - {name2} is significantly better")
else:
    print(f" Result: NOT statistically significant (p ≥ 0.05)")
    print(f" - No significant difference between models")
```

Logistic Regression vs Decision Tree:

```
Mean Difference: -0.0169
t-statistic: -1.4389
p-value: 0.1881
Result: NOT statistically significant (p ≥ 0.05)
- No significant difference between models
```

In [25]:

```
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

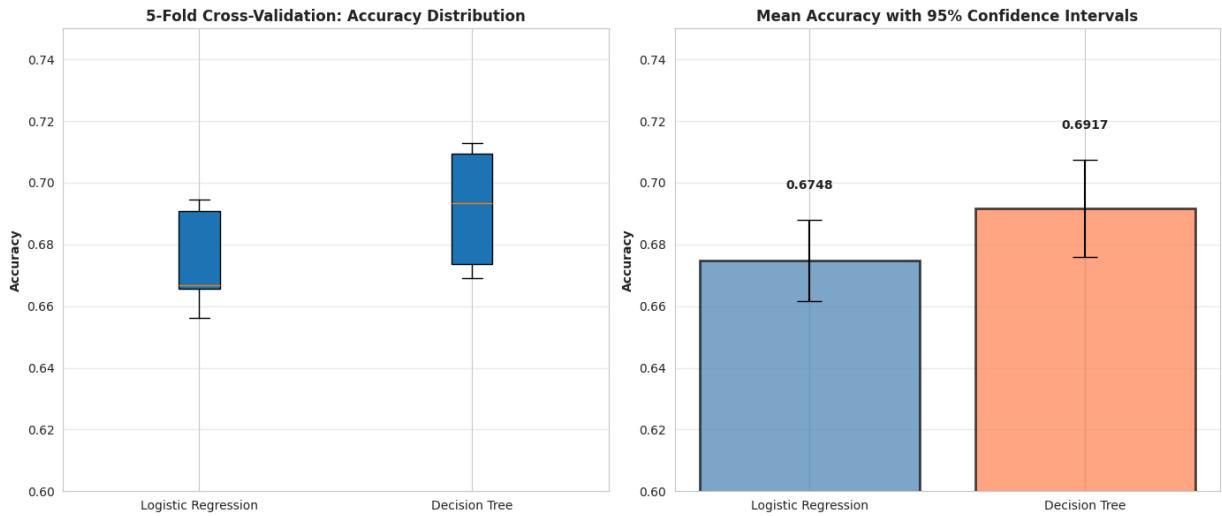
cv_data = [cv_results[name] for name in model_names]
axes[0].boxplot(cv_data, labels=model_names, patch_artist=True)
axes[0].set_ylabel('Accuracy', fontweight='bold')
axes[0].set_title('5-Fold Cross-Validation: Accuracy Distribution', fontweight='bold')
axes[0].set_ylim([0.6, 0.75])
axes[0].grid(True, alpha=0.3, axis='y')

means = [scores.mean() for scores in cv_data]
stds = [scores.std() for scores in cv_data]
ci_95 = [1.96 * std / np.sqrt(5) for std in stds]

axes[1].bar(model_names, means, yerr=ci_95, capsize=10, alpha=0.7,
            color=['steelblue', 'coral'], edgecolor='black', linewidth=2)
axes[1].set_ylabel('Accuracy', fontweight='bold')
axes[1].set_title('Mean Accuracy with 95% Confidence Intervals', fontweight='bold')
axes[1].set_ylim([0.6, 0.75])
axes[1].grid(True, alpha=0.3, axis='y')

for i, (mean, ci) in enumerate(zip(means, ci_95)):
    axes[1].text(i, mean + 0.01, f'{mean:.4f}', ha='center', fontweight='bold')

plt.tight_layout()
plt.show()
```



Causal Bias Decomposition Analysis

This analysis separates the sources of algorithmic bias into two components:

1. **Data Bias:** Pre-existing disparities in the training data
2. **Algorithmic Amplification:** Additional bias introduced by the machine learning model

Understanding this decomposition helps identify whether bias mitigation should focus on data collection or model design.

In [26]:

```
# Decomposing bias into data-driven vs algorithmic components
aa_recidivism_rate = y_train[sens_train == 1].mean()
other_recidivism_rate = y_train[sens_train == 0].mean()
baseline_data_di = aa_recidivism_rate / other_recidivism_rate

print(f"\nActual Recidivism Rates in Training Data:")
print(f"  African American: {aa_recidivism_rate:.4f}")
print(f"  Other Races: {other_recidivism_rate:.4f}")
print(f"  Disparate Impact: {baseline_data_di:.4f}")
print(f"  - Data already contains {(baseline_data_di - 1) * 100:.1f}% more positive")

# Comparing data bias to model bias
model_di = baseline_di
print(f"\nModel's Disparate Impact (on test set): {model_di:.4f}")

# Calculating how much the model amplifies existing bias
amplification_ratio = model_di / baseline_data_di
print(f"\nAmplification Factor: {amplification_ratio:.4f}")
print(f"  - Model amplifies data bias by {((amplification_ratio - 1) * 100:.1f)%}")

if amplification_ratio > 1:
    print(f"  - Model worsens bias")
elif amplification_ratio < 1:
    print(f"  - Model weakens bias")
else:
    print(f"  - Model preserves bias")

# Estimating relative contributions of data vs algorithm
data_bias_contribution = (baseline_data_di - 1) / (model_di - 1) * 100 if model_di >
algo_bias_contribution = 100 - data_bias_contribution

print(f"\nEstimated Contribution to Model Disparate Impact:")
print(f"  Data Bias: {max(0, data_bias_contribution):.1f}%")
print(f"  Algorithmic Amplification: {max(0, algo_bias_contribution):.1f}%")
```

```

print(f"\nInterpretation:")
if data_bias_contribution > 80:
    print(f" - Bias is primarily DATA-DRIVEN")
    print(f" - Solution: Rebalance data, fix upstream processes")
else:
    print(f" - Bias is primarily ALGORITHMIC")
    print(f" - Solution: Change model design, apply fairness constraints")

```

Actual Recidivism Rates in Training Data:

- African American: 0.5183
- Other Races: 0.3884
- Disparate Impact: 1.3344
- Data already contains 33.4% more positive cases for AA

Model's Disparate Impact (on test set): 2.1379

Amplification Factor: 1.6022

- Model amplifies data bias by 60.2%
- Model worsens bias

Estimated Contribution to Model Disparate Impact:

- Data Bias: 29.4%
- Algorithmic Amplification: 70.6%

Interpretation:

- Bias is primarily ALGORITHMIC
- Solution: Change model design, apply fairness constraints

In [27]:

```

# Visualizing bias decomposition and amplification effects
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

if data_bias_contribution > 0 and algo_bias_contribution > 0:
    sizes = [max(0, data_bias_contribution), max(0, algo_bias_contribution)]
    labels = [f'Data Bias\n({max(0, data_bias_contribution):.1f}%)',
              f'Algorithmic Amplification\n({max(0, algo_bias_contribution):.1f}%)']
    colors = ['#ff9999', '#ffcc99']
    explode = (0.05, 0.05)

    axes[0].pie(sizes, explode=explode, labels=labels, colors=colors, autopct='%1.1f%%',
                shadow=True, startangle=90, textprops={'fontsize': 11, 'fontweight': 'bold'})
    axes[0].set_title('Disparate Impact: Source Attribution', fontweight='bold', fontstyle='italic')

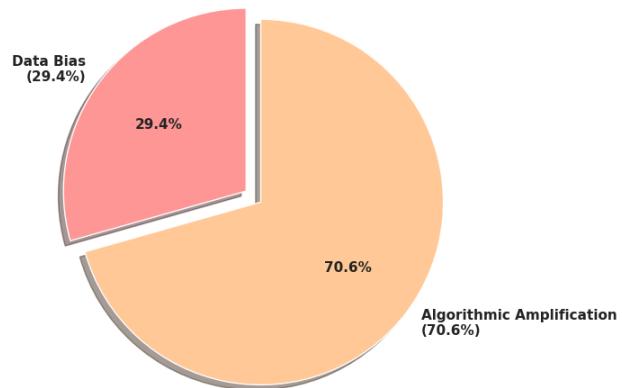
    stages = ['Data\n', 'Model\nPredictions']
    dis_values = [baseline_data_di, model_di]
    colors_bar = ['orange', 'red']

    axes[1].bar(stages, dis_values, color=colors_bar, edgecolor='black', linewidth=2, align='center')
    axes[1].axhline(y=1.0, color='green', linestyle='--', linewidth=2, label='Perfect Fairness')
    axes[1].set_ylabel('Disparate Impact Ratio', fontweight='bold', fontsize=11)
    axes[1].set_title('Disparate Impact: From Data to Model Predictions', fontweight='bold', fontstyle='italic')
    axes[1].set_ylim([0, max(dis_values) + 0.3])
    axes[1].legend(fontsize=10)
    axes[1].grid(True, alpha=0.3, axis='y')

for i, v in enumerate(dis_values):
    axes[1].text(i, v + 0.05, f'{v:.3f}', ha='center', fontweight='bold', fontsize=11)

plt.tight_layout()
plt.show()

```

Disparate Impact: Source Attribution**Disparate Impact: From Data to Model Predictions**